

Assignment2

February 19, 2022

Complete each problem below and print to pdf. Submit the pdf.
You will need to work with the three datasets attached to this assignment:

- poverty.csv
- poverty_2.csv
- real_estate.csv

1 Problem 1: Univariate Linear Regression

```
[123]: #export pdf
%%capture
!wget -nc https://raw.githubusercontent.com/brpy/colab-pdf/master/colab_pdf.py
from colab_pdf import colab_pdf
colab_pdf('pandas-assignment.ipynb')
```

```
ValueError                                Traceback (most recent call
last)

<ipython-input-123-a77988038c96> in <module>()
      1 get_ipython().system('wget -nc https://raw.githubusercontent.com/
brpy/colab-pdf/master/colab_pdf.py')
      2 from colab_pdf import colab_pdf
----> 3 colab_pdf('pandas-assignment.ipynb')

/content/colab_pdf.py in colab_pdf(file_name, notebookpath)
     20     # Check if the notebook exists in the Drive.
     21     if not os.path.isfile(os.path.join(notebookpath, file_name)):
--> 22         raise ValueError(f"file '{file_name}' not found in path
'{notebookpath}'.")
     23
     24     # Installing all the recommended packages.
```

```
ValueError: file 'pandas-assignment.ipynb' not found in path '/content/  
↳drive/MyDrive/Colab Notebooks/'.
```

1.1 1) import the libraries you will need:

numpy pandas matplotlib.pyplot statsmodels.api

```
[93]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import statsmodels.api as sm  
from google.colab import files
```

1.2 2) Import the date poverty.csv dataset

```
[94]: url = 'https://raw.githubusercontent.com/nnguyen09/Machine-Learning-/master/  
↳assignment2/poverty.csv'  
df = pd.read_csv(url)  
df.head()
```

```
[94]:
```

	Location	PovPct	Brth15to17	Brth18to19	ViolCrime	TeenBrth
0	Alabama	20.1	31.5	88.7	11.2	54.5
1	Alaska	7.1	18.9	73.7	9.1	39.5
2	Arizona	16.1	35.0	102.5	10.4	61.2
3	Arkansas	14.9	31.6	101.7	10.4	59.9
4	California	16.7	22.6	69.1	11.2	41.1

1.3 3) Print the dataset indexed upon the location column.

```
[95]: df['Location']
```

```
[95]:
```

0	Alabama
1	Alaska
2	Arizona
3	Arkansas
4	California
5	Colorado
6	Connecticut
7	Delaware
8	District_of_Columbia
9	Florida
10	Georgia
11	Hawaii
12	Idaho
13	Illinois

```

14         Indiana
15         Iowa
16         Kansas
17         Kentucky
18         Louisiana
19         Maine
20         Maryland
21         Massachusetts
22         Michigan
23         Minnesota
24         Mississippi
25         Missouri
26         Montana
27         Nebraska
28         Nevada
29         New_Hampshire
30         New_Jersey
31         New_Mexico
32         New_York
33         North_Carolina
34         North_Dakota
35         Ohio
36         Oklahoma
37         Oregon
38         Pennsylvania
39         Rhode_Island
40         South_Carolina
41         South_Dakota
42         Tennessee
43         Texas
44         Utah
45         Vermont
46         Virginia
47         Washington
48         West_Virginia
49         Wisconsin
50         Wyoming
Name: Location, dtype: object

```

1.4 4) Get useful descriptive statistical data on the dataset.

Hint: this is a single line, data._____

```
[96]: df.describe()
```

```

[96]:      PovPct  Brth15to17  Brth18to19  ViolCrime  TeenBrth
count  51.000000    51.000000    51.000000    51.000000    51.000000
mean   13.117647    22.282353    72.019608     7.854902    42.243137

```

std	4.277228	8.043499	18.975563	8.914131	12.318511
min	5.300000	8.100000	39.000000	0.900000	20.000000
25%	10.250000	17.250000	58.300000	3.900000	33.900000
50%	12.200000	20.000000	69.400000	6.300000	39.500000
75%	15.800000	28.100000	87.950000	9.450000	52.600000
max	25.300000	44.800000	104.300000	65.000000	69.100000

1.5 5) Print the columns

```
[97]: df.columns
```

```
[97]: Index(['Location', 'PovPct', 'Brth15to17', 'Brth18to19', 'ViolCrime',  
          'TeenBrth'],  
          dtype='object')
```

1.6 6) Create a regression line based upon the dependent and independent variables:

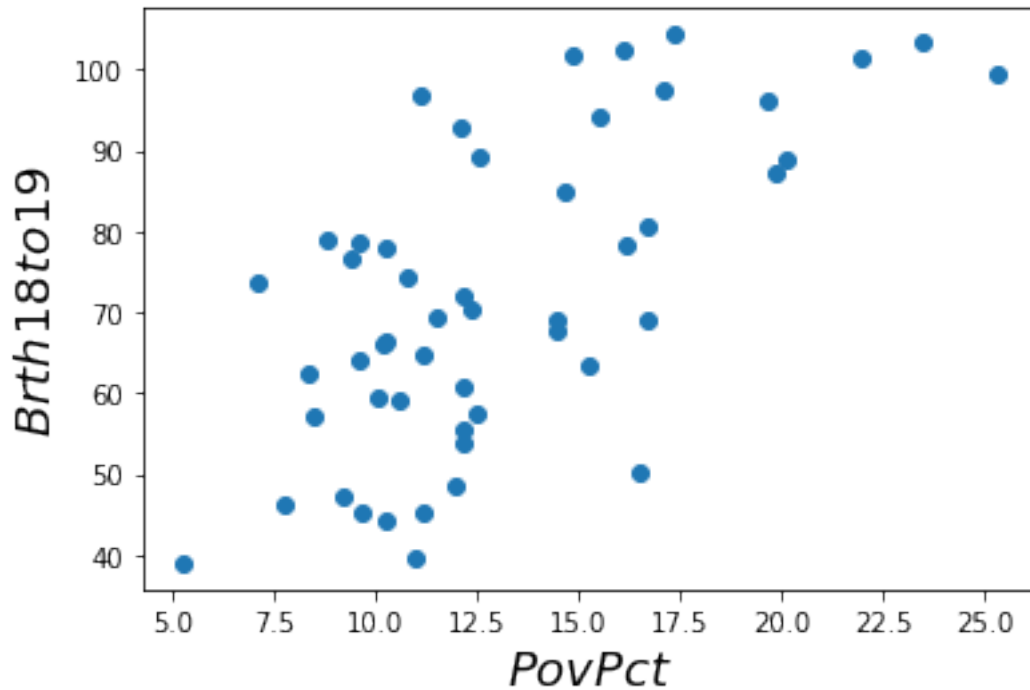
PovPct Brth18to19

In this step only create a scatterplot of the two variables, simply plotting the data.

Note: The variable PovPct is the percent of a state's population in 2000 living in households with incomes below the federally defined poverty level.

```
[98]: x=df['PovPct']  
      y=df['Brth18to19']  
      plt.xlabel("$PovPct$", fontsize = 18)  
      plt.ylabel("$Brth18to19$", fontsize = 18)  
      plt.scatter(x, y)
```

```
[98]: <matplotlib.collections.PathCollection at 0x7f97e7465f90>
```



1.7 7) Lets create a new variable, x1, as well as the results variable:

Example would be 1. `x1 = sm.add_constant(x)` 2. `results = sm.OLS(y, x1).fit()` 3. `results.summary()`

This gives you the OLS Regression results, the coefficients table, and some additional tests. The data that you are interested in is the coefficient values. This is the value for the constant you created is `b0`, and `birth19to19` is `b1` in the regression equation.

```
[99]: x1 = sm.add_constant(x)
      result = sm.OLS(y,x1).fit()
      print(result.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Brth18to19    R-squared:                  0.422
Model:                            OLS        Adj. R-squared:             0.410
Method:                 Least Squares    F-statistic:                 35.78
Date:                  Sat, 19 Feb 2022    Prob (F-statistic):          2.50e-07
Time:                      01:14:44    Log-Likelihood:              -207.98
No. Observations:                  51    AIC:                        420.0
Df Residuals:                      49    BIC:                        423.8
Df Model:                            1
Covariance Type:                  nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
=====
```

```

-----
const          34.2124      6.641      5.151      0.000      20.866      47.559
PovPct          2.8822      0.482      5.982      0.000      1.914      3.850
=====
Omnibus:                1.175      Durbin-Watson:                2.161
Prob(Omnibus):          0.556      Jarque-Bera (JB):          0.988
Skew:                   0.088      Prob(JB):                  0.610
Kurtosis:               2.341      Cond. No.                  45.1
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/tsatools.py:117:

FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

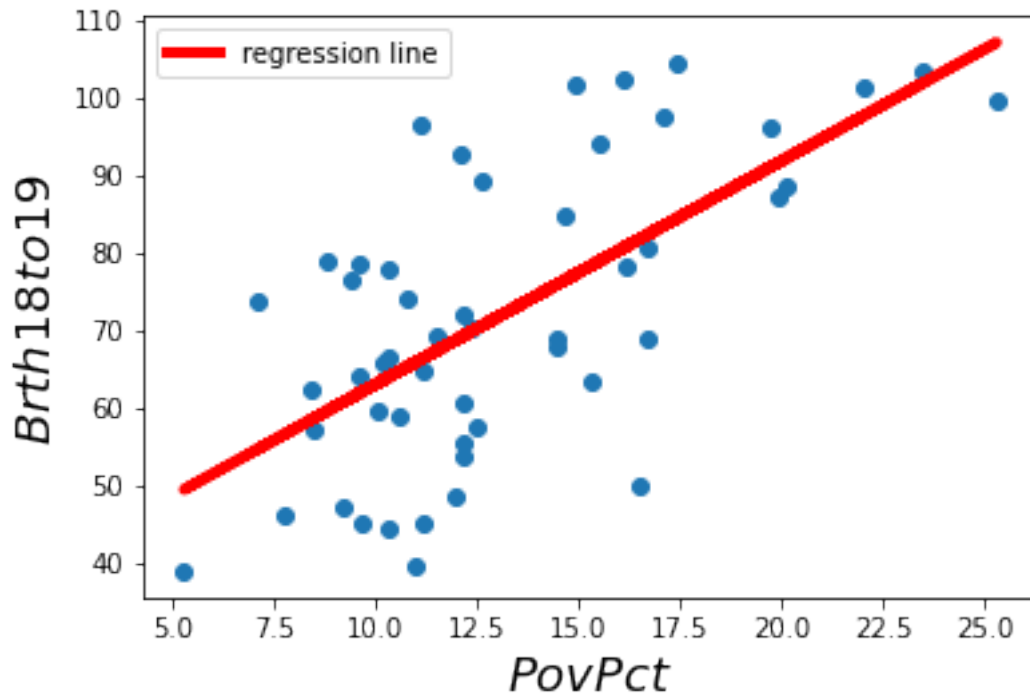
```
x = pd.concat(x[:,order], 1)
```

1.8 8) Taking the coefficient values for the new constant and the Y variable, create a scatterplot:

e.g. $\text{yhat} = 0.1464 \cdot x + 0.25712$ `fig = plt.plot(x, yhat, lw=4, c='red', label = 'regression line')`

```
[100]: yhat = 2.8822*x +34.2124
plt.xlabel("$PovPct$", fontsize = 18)
plt.ylabel("$Brth18to19$", fontsize = 18)
plt.scatter(x, y)
fig = plt.plot(x, yhat,lw=4, c= 'red',label = 'regression line')
plt.legend(loc='upper left')
```

```
[100]: <matplotlib.legend.Legend at 0x7f97e7465ed0>
```




2 Problem 2: Implement code from lecture

2.1 1) Perform linear regression using the normal equation, as done in slides.

```
[101]: plt.xlabel("$PovPct$", fontsize = 18)
plt.ylabel("$Brth18to19$", fontsize = 18)
plt.scatter(x,y)
X_b = np.c_[np.ones((x.size,1)),x] #add x0 = 1 to each instance
theta_best = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y)
theta_best
X_new= np.array([[0],[25]])
X_new_b = np.c_[np.ones((2,1)), X_new] #add x0 = 1 to each instance
y_predict = X_new_b.dot(theta_best)
y_predict
plt.plot(X_new, y_predict, c="red", label = "Prediction")
plt.legend(loc = "upper left")
```

```
[101]: <matplotlib.legend.Legend at 0x7f97e7386390>
```




2.2 2) Perform linear regression using Scikit-Learn, as done in the slides.

```
[102]: from sklearn.linear_model import LinearRegression
plt.scatter(x,y)
x_new = np.c_[x]
y_new = np.c_[y]

lin_reg = LinearRegression()
lin_reg.fit(x_new,y_new)
lin_reg.intercept_, lin_reg.coef_

plt.plot(x_new, lin_reg.coef_*x_new + lin_reg.intercept_, c='red',
        ↪label='Prediction')
plt.legend(loc = 'upper left')
```

[102]: <matplotlib.legend.Legend at 0x7f97e7350350>



3 Problem 3: Multivariate Linear Regression

In this problem we will continue using the poverty dataset. Do poverty and violent crimes affect teen pregnancy?

3.1 1) import the libraries you will need:

numpy pandas matplotlib.pyplot statsmodels.api

```
[103]: import numpy as np
import pandas as pd
from sklearn.preprocessing import normalize
```

3.2 2) Import the dataset, poverty_2.csv, and print it.

```
[104]: url_2 = 'https://raw.githubusercontent.com/nnguyen09/Machine-Learning-/master/
→assignment2/poverty_2.csv'
data= pd.read_csv(url_2)
print(data)
```

	PovPct	ViolCrime	TeenBrth
0	20.1	11.2	54.5
1	7.1	9.1	39.5
2	16.1	10.4	61.2
3	14.9	10.4	59.9

4	16.7	11.2	41.1
5	8.8	5.8	47.0
6	9.7	4.6	25.8
7	10.3	3.5	46.3
8	22.0	65.0	69.1
9	16.2	7.3	44.5
10	12.1	9.5	55.7
11	10.3	4.7	38.2
12	14.5	4.1	39.1
13	12.4	10.3	42.2
14	9.6	8.0	44.6
15	12.2	1.8	32.5
16	10.8	6.2	43.0
17	14.7	7.2	51.0
18	19.7	17.0	58.1
19	11.2	2.0	25.4
20	10.1	11.8	35.4
21	11.0	3.6	23.3
22	12.2	8.5	34.8
23	9.2	3.9	27.5
24	23.5	12.9	64.7
25	9.4	8.8	44.1
26	15.3	3.0	36.4
27	9.6	2.9	37.0
28	11.1	10.7	53.9
29	5.3	1.8	20.0
30	7.8	5.1	26.8
31	25.3	8.8	62.4
32	16.5	8.5	29.5
33	12.6	9.4	52.2
34	12.0	0.9	27.2
35	11.5	5.4	39.5
36	17.1	12.2	58.0
37	11.2	4.1	36.8
38	12.2	6.3	31.6
39	10.6	3.3	35.6
40	19.9	7.9	53.0
41	14.5	1.8	38.0
42	15.5	10.6	54.3
43	17.4	9.0	64.4
44	8.4	3.9	36.8
45	10.3	2.2	24.2
46	10.2	7.6	37.6
47	12.5	5.1	33.0
48	16.7	4.9	45.5
49	8.5	4.3	32.3
50	12.2	2.1	39.9

3.3 3) We need to normalize the input variables.

```
[105]: from sklearn.preprocessing import normalize
data = normalize(data, axis=0)
```

3.4 4) Split the data into input variables, X, and the output variable, Y.

```
[106]: X=data[:,0:2]
Y=data[:,2:]
```

3.5 5) Graph the dataset with a seed of 42.

Replace the FILLINTHESEVALUES fields.

```
[107]: np.random.seed(42)


fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')

xs = 2*np.random.rand(100,1)
ys = 2+3*xs+np.random.rand(100,1)
zs = 4*xs+3*ys+np.random.rand(100,1)

ax.scatter(xs, ys, zs)

ax.set_xlabel('X')
ax.set_ylabel('Y')
ax.set_zlabel('Z')

plt.show()
```



3.6 6) Implement Gradient Descent.

This section has been provided. Please run and understand the code.

```
[108]: # hyperparameters
learning_rate = 0.05
max_iteration = 500

#parameters
theta = np.zeros((data.shape[1], 1))

[109]: def hypothesis (theta, X) :
    tempX = np.ones((X.shape[0], X.shape[1] + 1))
    tempX[:,1:] = X
    return np.matmul(tempX, theta)

[110]: def loss (theta, X, Y) :
    return np.average(np.square(Y - hypothesis(theta, X))) / 2

[111]: def gradient (theta, X, Y) :
    tempX = np.ones((X.shape[0], X.shape[1] + 1))
    tempX[:,1:] = X
    d_theta = - np.average((Y - hypothesis(theta, X)) * tempX, axis= 0)
    d_theta = d_theta.reshape((d_theta.shape[0], 1))
    return d_theta

[112]: def gradient_descent (theta, X, Y, learning_rate, max_iteration, gap) :
    cost = np.zeros(max_iteration)
    for i in range(max_iteration) :
```

```

d_theta = gradient (theta, X, Y)
theta = theta - learning_rate * d_theta
cost[i] = loss(theta, X, Y)
if i % gap == 0 :
    print ('iteration : ', i, ' loss : ', loss(theta, X, Y))
return theta, cost

```

```

[113]: # Training model
theta, cost = gradient_descent (theta, X, Y, learning_rate, max_iteration, 100)

```

```

iteration : 0  loss : 0.008893757788504215
iteration : 100  loss : 0.0006811106575134702
iteration : 200  loss : 0.0006573219302696655
iteration : 300  loss : 0.0006360731168287809
iteration : 400  loss : 0.0006169026951758099

```

```

[114]: #optimal value is :
theta

```

```

[114]: array([[0.12381477],
              [0.04264512],
              [0.05698502]])


```

```

[115]: #plot cost
fig, ax = plt.subplots()
ax.plot(np.arange(max_iteration), cost, 'r')
ax.legend(loc='upper right', labels=['batch gradient descent'])
ax.set_xlabel('Iterations')
ax.set_ylabel('Cost')
ax.set_title('Error vs. Training Epoch')

plt.show()

```



3.7 7) Implement Stochastic Gradient Descent. Please run.

```
[116]: def stochastic_gradient_descent (theta, X, Y, learning_rate, max_iteration,
    ↪gap) :
    cost = np.zeros(max_iteration)
    for i in range(max_iteration) :
        for j in range(X.shape[0]):
            d_theta = gradient (theta, X[j,:].reshape(1, X.shape[1]), Y[j,:].
    ↪reshape(1, 1))
            theta = theta - learning_rate * d_theta

        cost[i] = loss(theta, X, Y)
        if i % gap == 0 :
            print ('iteration : ', i, ' loss : ', loss(theta, X, Y))
    return theta, cost
```

```
[53]: theta_stoc = np.zeros((data.shape[1], 1))
```


```
[54]: theta_stoc, cost_stoc = stochastic_gradient_descent (theta_stoc, X, Y,
    ↪learning_rate, max_iteration, 100)
```

```
iteration : 0 loss : 0.0007764556902156442
iteration : 100 loss : 0.0004037848207345314
iteration : 200 loss : 0.00036553095210465356
```

```
iteration : 300 loss : 0.000347847758744226
iteration : 400 loss : 0.00033956148785195
```

```
[117]: #plot the cost
fig, ax = plt.subplots()
ax.plot(np.arange(max_iteration), cost, 'r')
ax.plot(np.arange(max_iteration), cost_stoc, 'b')
#ax.plot(np.arange(max_iteration), mb_cost, 'g')
ax.legend(loc='upper right', labels=['batch gradient descent', 'stochastic_
→gradient descent'])#, 'mini-batch gradient descent'])
ax.set_xlabel('Iterations')
ax.set_ylabel('Cost')
ax.set_title('Error vs. Training Epoch')

plt.show()
```



```
[118]: np.random.seed(42)

fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')

xs = X[:, 0]
ys = X[:, 1]
zs = Y
```

```

ax.scatter(xs, ys, zs)


ax.set_xlabel('PovPct')
ax.set_ylabel('ViolCrime')
ax.set_zlabel('TeenBrth')

# new
x = y = np.arange(0, 0.3, 0.05)
xp, yp = np.meshgrid(x, y)
z = np.array([hypothesis(theta, np.array([x,y]))[0, 0] for x,y in zip(np.
    →ravel(xp), np.ravel(yp))])
zp = z.reshape(xp.shape)

ax.plot_surface(xp, yp, zp, alpha=0.7)

plt.show()

```



4 Problem 4, predict house price.

- import real_estate.csv
- Are there any null values in the dataset? Drop any missing data if exist.
- Create X as a 1-D array of the distance to the nearest MRT station, and y as the housing price
- What is the number of samples in the data set? To do this, you can look at the "shape" of X and y
- Split the data into train and test sets using sklearn's train_test_split, with test_size = 1/3

- Find the line of best fit using a Linear Regression and show the result of coefficients and intercept (you can use sklearn's linear regression)
- Using the predict method, make predictions for the test set and evaluate the performance (e.g., MSE or other metrics).

```
[119]: #import real_estate.csv
url_3 = 'https://raw.githubusercontent.com/nnguyen09/Machine-Learning-/master/
→assignment2/real_estate.csv'
data_real_estate = pd.read_csv(url_3)

print(data_real_estate.columns)

#Are there any null values in the dataset? Drop any missing data if exist.
data_real_estate.dropna(axis = 0, how = 'any')

#Create X as a 1-D array of the distance to the nearest MRT station, and y as
→the housing price
X=data_real_estate['X3 distance to the nearest MRT station'].array
print(X)
Y=data_real_estate['Y house price of unit area'].array
print(Y)

#What is the number of samples in the data set? To do this, you can look at the
→"shape" of X and y
print('Number of sample in data set is: ' , X.shape, Y.shape)
```

```
Index(['No', 'X1 transaction date', 'X2 house age',
      'X3 distance to the nearest MRT station',
      'X4 number of convenience stores', 'X5 latitude', 'X6 longitude',
      'Y house price of unit area'],
      dtype='object')
<PandasArray>
[84.87882, 306.5947, 561.9845, 561.9845, 390.5684, 2175.03, 623.4731,
 287.6025, 5512.038, 1783.18,
 ...
 289.3248, 130.9945, 372.1386, 2408.993, 2175.744, 4082.015, 90.45606,
 390.9696, 104.8101, 90.45606]
Length: 414, dtype: float64
<PandasArray>
[37.9, 42.2, 47.3, 54.8, 43.1, 32.1, 40.3, 46.7, 18.8, 22.1,
 ...
 41.2, 37.2, 40.5, 22.3, 28.1, 15.4, 50.0, 40.6, 52.5, 63.9]
Length: 414, dtype: float64
Number of sample in data set is: (414,) (414,)
```

```
[120]: #Split the data into train and test sets using sklearn's train_test_split, with
→test_size = 1/3
```

```

from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

X_train, X_test, y_train, y_test = train_test_split(X,Y,test_size = 1/3)

X_train = X_train.reshape(-1, 1)
y_train = y_train.reshape(-1, 1)
X_test = X_test.reshape(-1,1)
y_test = y_test.reshape(-1,1)

```

[121]: *#Find the line of best fit using a Linear Regression and show the result of*
→coefficients and intercept (you can use sklearn's linear regression)

```

lin_reg = LinearRegression()
lin_reg.fit(X_train,y_train)
print(lin_reg.intercept_, lin_reg.coef_)

```

```
[46.68201581] [[-0.00764843]]
```

[122]: *#Using the predict method, make predictions for the test set and evaluate the*
→performance (e.g., MSE or other metrics).

```

y_predict = lin_reg.predict(X_test)

print(mean_squared_error(y_test, y_predict))
print(lin_reg.score(X_test,y_test))

```

```

91.52078612852627
0.42585247237704293

```