

COMP4220: Machine Learning, Spring 2022, Assignment 3

Please submit one pdf file for all questions.

#importing the libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from google.colab import files
```

```
url =
'https://raw.githubusercontent.com/nnguyen09/Machine-Learning-/master/
assignment3/wine.csv'
df = pd.read_csv(url)
df
```

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
0	7.4	0.700	0.00	1.9
0.076				
1	7.8	0.880	0.00	2.6
0.098				
2	7.8	0.760	0.04	2.3
0.092				
3	11.2	0.280	0.56	1.9
0.075				
4	7.4	0.700	0.00	1.9
0.076				
...
...				
1594	6.2	0.600	0.08	2.0
0.090				
1595	5.9	0.550	0.10	2.2
0.062				
1596	6.3	0.510	0.13	2.3
0.076				
1597	5.9	0.645	0.12	2.0
0.075				
1598	6.0	0.310	0.47	3.6
0.067				

	free sulfur dioxide	total sulfur dioxide	density	pH
sulphates \				
0	11.0	34.0	0.99780	3.51
0.56				
1	25.0	67.0	0.99680	3.20
0.68				

2	15.0	54.0	0.99700	3.26
0.65				
3	17.0	60.0	0.99800	3.16
0.58				
4	11.0	34.0	0.99780	3.51
0.56				
...
...				
1594	32.0	44.0	0.99490	3.45
0.58				
1595	39.0	51.0	0.99512	3.52
0.76				
1596	29.0	40.0	0.99574	3.42
0.75				
1597	32.0	44.0	0.99547	3.57
0.71				
1598	18.0	42.0	0.99549	3.39
0.66				

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

[1599 rows x 12 columns]

variables (based on physicochemical tests):

Tips

An interesting thing to do is to set an arbitrary cutoff for your dependent variable (wine quality): 7 or higher getting classified as '1' and the remainder as '0'.


This allows you to convert this problem into a classification problem.

1. Since we want to classify the wine base on the quality so we want to look at the distribution of the wine quality

Make a histogram plot for the quality column to see the distribution of the wine quality

```
x='quality'  
df.hist(column=x)
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at  
0x7efddb9040d0>]],  
      dtype=object)
```



2. Show the number of null values using sum() method. If there are null values then remove them from the dataset

```
df.isnull().sum()
```

fixed acidity	0
volatile acidity	0
citric acid	0
residual sugar	0
chlorides	0
free sulfur dioxide	0
total sulfur dioxide	0

```

density          0
pH               0
sulphates        0
alcohol          0
quality          0
dtype: int64

```

3. Since we want to categorize the dependent variable (wine quality)

Change the quality column to 1 if the quality ≥ 7 , and 0 if the quality is < 7

Show the dataset after making this change

Hint: the quality column should only have 0s and 1s after the change

```

df['quality'] = [1 if x >= 7 else 0 for x in df['quality']]
df

```

```

      fixed acidity  volatile acidity  citric acid  residual sugar
chlorides \
0          7.4          0.700          0.00          1.9
0.076
1          7.8          0.880          0.00          2.6
0.098
2          7.8          0.760          0.04          2.3
0.092
3         11.2          0.280          0.56          1.9
0.075
4          7.4          0.700          0.00          1.9
0.076
...          ...          ...          ...          ...
...
1594         6.2          0.600          0.08          2.0
0.090
1595         5.9          0.550          0.10          2.2
0.062
1596         6.3          0.510          0.13          2.3
0.076
1597         5.9          0.645          0.12          2.0
0.075
1598         6.0          0.310          0.47          3.6
0.067

```

```

      free sulfur dioxide  total sulfur dioxide  density  pH
sulphates \
0          11.0          34.0  0.99780  3.51
0.56
1          25.0          67.0  0.99680  3.20
0.68
2          15.0          54.0  0.99700  3.26
0.65

```

3	17.0	60.0	0.99800	3.16
0.58				
4	11.0	34.0	0.99780	3.51
0.56				
...
...				
1594	32.0	44.0	0.99490	3.45
0.58				
1595	39.0	51.0	0.99512	3.52
0.76				
1596	29.0	40.0	0.99574	3.42
0.75				
1597	32.0	44.0	0.99547	3.57
0.71				
1598	18.0	42.0	0.99549	3.39
0.66				

	alcohol	quality
0	9.4	0
1	9.8	0
2	9.8	0
3	9.8	0
4	9.4	0
...
1594	10.5	0
1595	11.2	0
1596	11.0	0
1597	10.2	0
1598	11.0	0

[1599 rows x 12 columns]

4. Create y as the quality column and X as everything but the quality column

```
# X=good quality, Y=bad quality
X = df.drop(['quality'], axis = 'columns')
y = df['quality']
df['quality'].value_counts()
```

```
0    1382
1     217
Name: quality, dtype: int64
```

5. Split the dataset into the training and test set using "train_test_split".

Split the training and test set into 70-30 ratio

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=.30, random_state=42)
```

6. Apply Feature Scaling method for X_train and X_test with "StandardScaler" from "sklearn.preprocessing"

Hint: use `StandardScaler.fit_transform` for "X_train" and use `StandardScaler.transform` for "X_test"

```
from sklearn.preprocessing import StandardScaler
```

```
scalar = StandardScaler()
```

```
X_train = scalar.fit_transform(X_train)
print(X_train)
```

```
X_test = scalar.transform(X_test)
```

```
[ [ 1.69536131e-01 -1.72107140e+00 4.59303345e-01 ... 1.01180685e+00
    1.22661179e+00 5.50057013e-01]
  [ 2.44606730e+00 -4.01957443e-01 1.84105501e+00 ... -2.10687612e+00
    1.22661179e+00 -2.05174641e-01]
  [-6.47680186e-01 3.77472102e-02 -1.28054303e-03 ... 4.92026353e-01
    2.97270776e-01 5.50057013e-01]
  ...
  [-6.47680186e-01 4.77451864e-01 -1.07597628e+00 ... 1.27169710e+00
    -6.90154049e-01 -8.66002338e-01]
  [-2.39072027e-01 -1.83099757e+00 4.08127357e-01 ... 3.72184202e-02
    8.20025095e-01 1.39969262e+00]
  [-1.46489650e+00 -1.33632983e+00 -5.24565306e-02 ... 4.92026353e-01
    -6.90154049e-01 2.91015593e+00]]
```

7. Train the logistic regression model on the training set using (solver='lbfgs', random_state = 42, max_iter = 1000)

```
from sklearn.linear_model import LogisticRegression
```

```
log_reg = LogisticRegression(solver='lbfgs', random_state = 42,  
max_iter = 1000)
```

```
log_reg.fit(X_train, y_train)
```

```
LogisticRegression(max_iter=1000, random_state=42)
```

8. Predict the results of x test

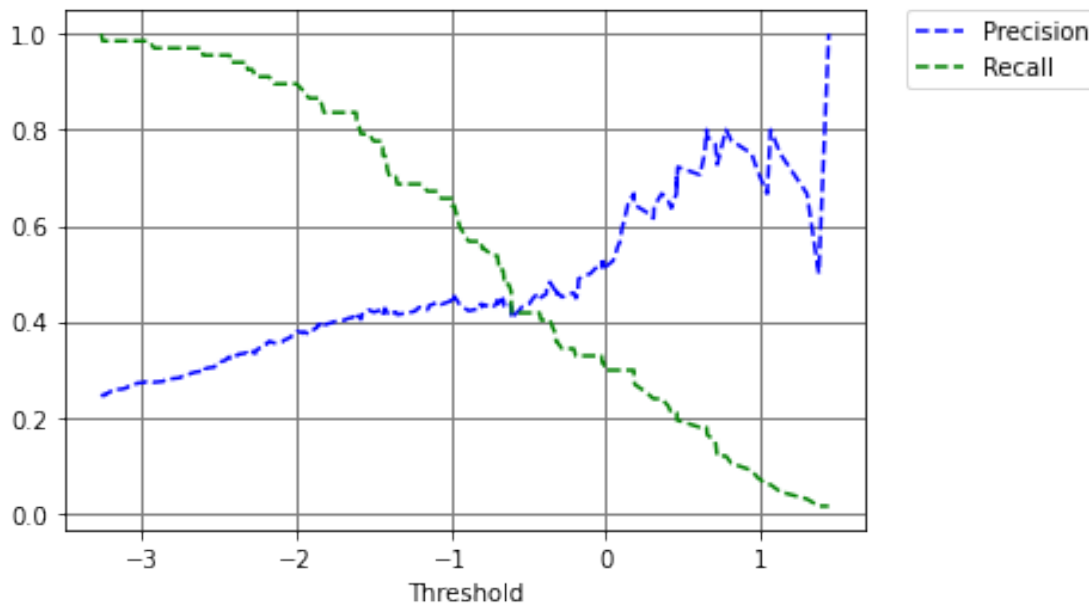
```
y_test_pred = log_reg.predict(X_test)
print(y_test_pred)
```

[illegible]

12. Use Matplotlib to plot precision and recall as functions of the threshold value

```
import matplotlib.pyplot as plt
def plot_precision_recall_vs_thresholds(precisions, recalls,
thresholds):
    plt.plot(thresholds, precisions[:-1], "b--", label="Precision")
    plt.plot(thresholds, recalls[:-1], "g--", label="Recall")
    plt.xlabel("Threshold")
    plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left',
borderaxespad=0.)
    plt.grid(b=True, which="both", axis="both", color='gray',
linestyle='-', linewidth=1)


plot_precision_recall_vs_thresholds(precisions, recalls, thresholds)
plt.show()
```



13. Plot the precision vs recall plot

```
fig, ax = plt.subplots()
ax.plot(recalls, precisions)
ax.set_ylabel('Precision')
ax.set_xlabel('Recall')
plt.show


<function matplotlib.pyplot.show>
```

14. Plot the ROC Curve

```
from sklearn.metrics import roc_curve
pr, tpr, thresh = roc_curve(y_test, y_score)
fig, ax = plt.subplots()
ax.plot(pr, tpr)
ax.set_ylabel('True Positive Rate(Recall)')
ax.set_xlabel('False Positive Rate')
plt.show
```

<function matplotlib.pyplot.show>



15. Find the area under the ROC Curve

```
from sklearn.metrics import roc_auc_score  
print (roc_auc_score(y_test, y_score))
```

0.8816450435473963